

# ECON7333: Assignment 3

Thomas Doyle

## Setup

```
options(scipen=999)
attach(ISLR::Wage)
```

## Exercise 1

### Exercise 1.1

The function `loocv_tmse` computes the test mean square error of a regression model using a leave-one-out cross-validation approach.

```
loocv_tmse <- function(d){
  #' @d a data.frame returned by `model.frame()`
  #
  n <- dim(d)[1]
  p <- dim(d)[2]

  MSE <- rep(0,n)

  for(i in 1:n) {
    #
    lm_i <- lm(d[-i,],y=TRUE) # leave i out
    MSE_i <- (lm_i$y[i]-predict(lm_i,d[i,]))^2
    MSE[i] <- MSE_i
  }

  return(
    #
    CV_n <- mean(MSE,na.rm = TRUE)
  )
}
```

### Exercise 1.2

We use `loocv_tmse` to choose between the three models;

1.  $\logwage = \beta_0 + \beta_1 age$
2.  $\logwage = \beta_0 + \beta_1 age + \beta_2 age^2$
3.  $\logwage = \beta_0 + \beta_1 age + \beta_2 education$

Each model takes `logwage` as the dependent variable.

```
Wage.models <- list(
  model.frame("logwage~age", ISLR::Wage),
  model.frame("logwage~age+I(age^2)", ISLR::Wage),
  model.frame("logwage~age+education", ISLR::Wage)
)
sapply(Wage.models, loocv_tmse)
```

```
## [1] 0.1306018 0.1381353 0.1531854
```

Choose the model with lowest test MSE. Choose linear regression on  $\logwage = \beta_0 + \beta_1 \text{age}$ , which returns the smallest test MSE.

### Exercise 1.3

The  $\lambda$  parameter is now used as tuning parameter on the ridge regression model;

$$\logwage = \beta_0 + \beta_1 \text{age} + \beta_2 \text{education}$$

```
lm.ridge <- function(lambda=0) {

  m <- model.frame(logwage~age+education, ISLR::Wage)
  Terms <- attr(eval.parent(m), "terms")
  Y <- model.response(m)
  X <- model.matrix(object = Terms, data = m, contrasts.arg = list(education="contr.treatment"))
  n <- nrow(X)
  p <- ncol(X)

  # https://arxiv.org/pdf/1509.09169.pdf
  # Hoerl, A. E. and Kennard, R. W. (1970).
  ridge.betas <- solve( t(X) %*% X + lambda*diag(p) ) %*% t(X) %*% Y

  l2.norm <- sqrt(sum(ridge.betas^2)) # drop intercept

  return(
    list(
      log.lambda=log(lambda),
      l2.norm=l2.norm,
      coefficients=ridge.betas,
      lambda=lambda
    )
  )
}
```

```
# MLE
# showing off
lm.ridge.loocv <- function(limits) {
  #' @limits
  #

  m <- model.frame(logwage~age+education, ISLR::Wage)
  Terms <- attr(eval.parent(m), "terms")
  Y <- model.response(m)
  X <- model.matrix(object = Terms, data = m, contrasts.arg = list(education="contr.treatment"))
```

```

loocv <- function(lambda, X, Y, Delta){
  n <- nrow(X)
  p <- ncol(X)
  loss <- 0

  for (i in 1:n) {
    loo_beta <- solve(t(X[-i,])%*%X[-i,]+lambda*diag(p))%*%t(X[-i,])%*%Y[-i]
    loss <- loss+(Y[i]-X[i,1]*loo_beta[1]-X[i,-1]%*%loo_beta[-1])^2
  }

  return(loss)
}

# optimize penalty parameter
# minimize RSS
opt <- optimize(loocv, limits, X=X, Y=Y)
l2.norm <- lm.ridge(opt$minimum)$l2.norm

return(
  data.frame(opt,l2.norm)
)
}

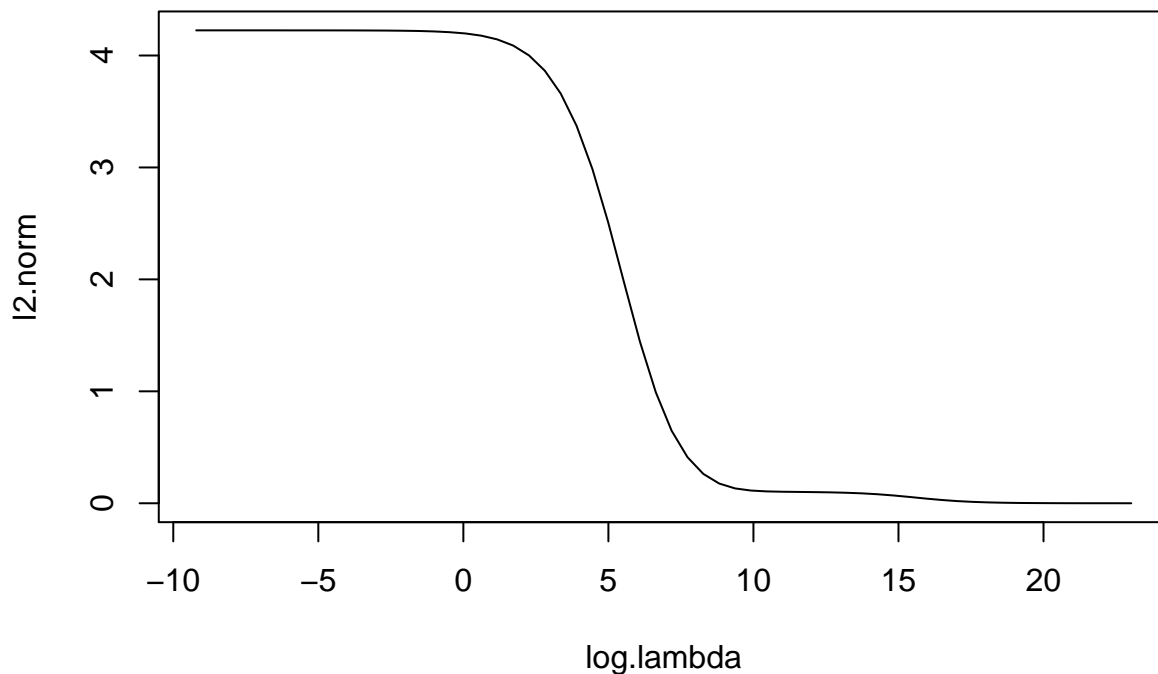
```

Plot  $\ln(\lambda)$  against  $\ell_2$  norm.

```

lambdas <- 100^seq(-2, 5, length = 60)
plot(t(sapply(lambdas,lm.ridge))[,1:2],type="l")

```



```

# mle
# lm.ridge.loocv(c(10^-10, 10^10))

```

## Exercise 2

### Exercise 2.1

The function `dgf` generates data following the data generating process set out by question 2 of ECON 7333 Assignment 3 text.

```
dgf <- function(n,p) {  
  #' @n `scalar`  
  #' @p `scalar`  
  
  e <- rnorm(n,0,1)  
  X <- matrix(runif(n*p),n,p)  
  y <- 2*X[,1]+4*X[,2]+e  
  
  output <- list(y=y,X=X)  
  return(list(y,X))  
}  
  
p.5 <- dgf(100,5)  
y <- p.5[[1]]  
X <- p.5[[2]]
```

### Exercise 2.2

Apply the model to the sample generated by `dgf`. `model_path` implements the subset selection algorithm set by the assignment text..

```
model_path <- function(y,X,M) {  
  #' @y `vector`, response variable  
  #' @X `matrix`, input variable  
  #' @M `scalar`, model path  
  
  p <- dim(X)[2]  
  n <- dim(X)[1]  
  
  # Initialise  
  #  
  b <- matrix(rep(NA,p*M),p,M)  
  b[,1] <- 0  
  r <- as.vector(y)  
  # e <- as.vector(rep(0.01,n))  
  e <- 0.01  
  
  for(m in 2:M) {  
    z <- b[,m-1]  
  
    ## select x_j with highest correlation with r  
    j <- which.max(cor(x=X,y=r))  
    xj <- X[,j]  
  
    a <- suppressWarnings(e*sign(t(xj)%*%r))  
    z[j] <- z[j]+a
```

```

    b[,m] <- z
    r <- as.vector(r-a%*%xj)
  }

  return(b)
}

```

### Exercise 2.3

Generate a sample path of length  $M = 1000$  of the coefficient estimates  $\hat{\beta}_j^{(m)}$ , and plot the sample using `matplot`.


```

path <- model_path(y,X,1000)

par(mfrow=c(1,2))
matplot(path,pch=2, cex=.2, col="black",
        ylab="level estimates",
        xlab=expression(hat(beta)[j]^(m)),
        main="")

plot(lm(y~X-1)$coefficients, pch=20, cex=.8,
     ylab="level coefficient estimates",
     xlab=expression(hat(beta)),
     main="y~X")

```



## Exercise 2.4

Divide the sample into training and test subsamples of size 50 each. Use cross-validation to find the best model along the paths (that is, find the best model among the  $M$  different models generated along the path). Choose between that best model and linear regression conducted by least squares.

```
mse_i <- function(y,X,m,i) {
  mean(sum((y[i] - m%*%t(X[i,]))^2))
}

subset_selection <- function(y,X,M,k) {
  #' @y
  #' @X
  #' @M
  #' @k

  # calculate MSE by iterate over M
  MSE_i <- apply(M,2,function(m) {
    mse_i(y,X,m,k)
  })

  # Return coef for model with min MSE
  return(
    list(
      coefficients = M[,which.min(MSE_i)],
      best_subset_m = which.min(MSE_i),
      best_mse = min(MSE_i)
    )
  )
}

k_fold <- function(n,k) {

  ix <- sample(1:n,n)

  fold_i <- split(ix,cut(1:n,k,FALSE))

  folds <- data.frame(fold_i)

  colnames(folds) <- paste("fold",1:k,sep = "_")

  return(folds)
}
```

The function `k_fold` returns an  $n \times k$  vector used to create  $k$  subsets of sample size  $n/k$  from the matrix of independent variables,  $X$ .

Use function `subset_selection` to select the best model from the training sample and compute the trainse MSE using `mse_1`.

Based on the minimum average MSE of  $k$  folds, choose the best model of coefficients. Compare the best subset model against the linear regression model.

```
folds <- k_fold(100,2)

# obtain test MSE using k-fold
```

```
#
fold_1_tmse = mse_i(y,X,subset_selection(y,X,path,folds[,1])[["coefficients"]],folds[,2])
fold_2_tmse = mse_i(y,X,subset_selection(y,X,path,folds[,2])[["coefficients"]],folds[,1])

SS_MSE <- subset_selection(y,X,path,folds[,which.min(c(fold_1_tmse,fold_2_tmse))])$best_mse
LM_MSE <- sum(lm(y~X-1)$residuals^2)

data.frame(SS_MSE,LM_MSE)
```

```
##      SS_MSE   LM_MSE
## 1 64.79127 117.5728
```

Choose best subset selection model, which returns the lowest CV k-fold MSE when compared to linear regression.

## Exercise 2.5

```
CV_k <- function(n,p) {
  # call dgf to generate new n x p data
  d <- dgf(n,p)

  y <- d[[1]]
  X <- d[[2]]

  # generate models
  path <- model_path(y,X,1000)

  folds <- k_fold(n,2)

  # obtain test MSE using k-fold
  tmse_1 <- mse_i(y,X,subset_selection(y,X,path,folds[,1])[["coefficients"]],folds[,2])
  tmse_2 <- mse_i(y,X,subset_selection(y,X,path,folds[,2])[["coefficients"]],folds[,1])

  SS_MSE <- subset_selection(y,X,path,folds[,which.min(c(fold_1_tmse,fold_2_tmse))])$best_mse
  LM_MSE <- sum(lm(y~X-1)$residuals^2)
  return(data.frame(SS_MSE,LM_MSE))
}
```

```
CV_k(100,5)
```

```
##      SS_MSE   LM_MSE
## 1 42.98299 96.15328
```

```
CV_k(100,10)
```

```
##      SS_MSE   LM_MSE
## 1 54.17434 97.33566
```

```
CV_k(100,20)
```

```
##      SS_MSE LM_MSE
## 1 50.23396 92.566
```

```
CV_k(100,50)
```

```
##      SS_MSE   LM_MSE  
## 1 47.56949 60.68574
```

As  $p$  increases, the linear regression MSE decreases while the MSE reported by best subset selection is relatively stable. At each level of  $p$ , the best subset selection MSE is lower than the linear regression MSE.